



Stochastic Gradient Descent

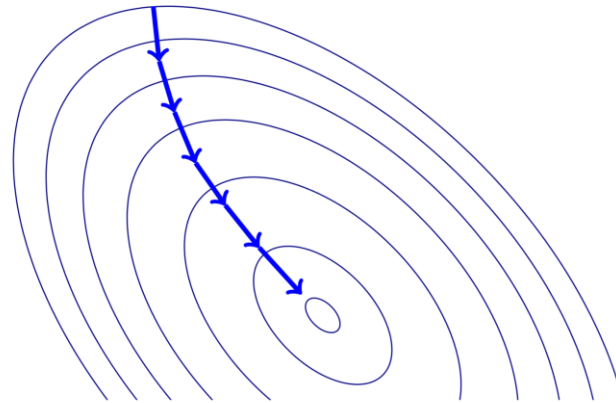
Prof. Seungchul Lee
Industrial AI Lab.

Gradient Descent

- We will cover gradient descent algorithm and its variants:
 - Batch Gradient Descent
 - Stochastic Gradient Descent
 - Mini-batch Gradient Descent
- We will explore the concept of these three gradient descent algorithms with a logistic regression model in TensorFlow
- Limitation of the Gradient Descent
 - Adaptive learning rate

Batch Gradient Descent (= Gradient Descent)

Repeat: $\omega \leftarrow \omega - \alpha \nabla f(\omega)$ for some step size (or learning rate) $\alpha > 0$



Batch Gradient Descent

- Loss function ℓ has been the average loss over all of the training examples:

$$\mathcal{E}(\omega) = \frac{1}{m} \sum_{i=1}^m \ell(\hat{y}_i, y_i) = \frac{1}{m} \sum_{i=1}^m \ell(h_{\omega}(x_i), y_i)$$

- By linearity,

$$\nabla_{\omega} \mathcal{E} = \nabla_{\omega} \frac{1}{m} \sum_{i=1}^m \ell(h_{\omega}(x_i), y_i) = \frac{1}{m} \sum_{i=1}^m \frac{\partial}{\partial \omega} \ell(h_{\omega}(x_i), y_i)$$

$$\omega \leftarrow \omega - \alpha \nabla_{\omega} \mathcal{E}$$

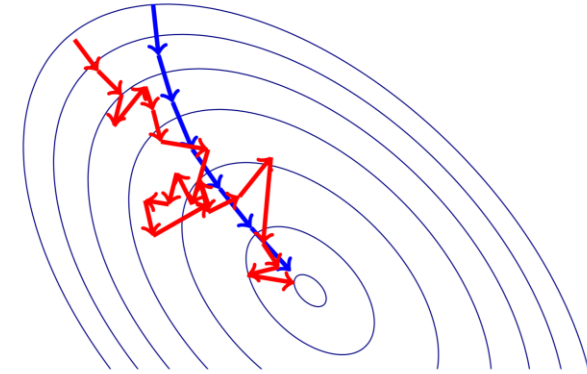
- Computing the gradient requires summing over **all of the training examples**.
- This is known as batch training.
- Batch training is **impractical** if you have a large dataset (e.g. millions of training examples) !

Stochastic Gradient Descent (SGD)

- Stochastic gradient descent (SGD): update the parameters based on the gradient for a randomly selected **single** training example:

$$\ell(\hat{y}_i, y_i) = \ell(h_{\omega}(x_i), y_i) = \ell^{(i)}$$

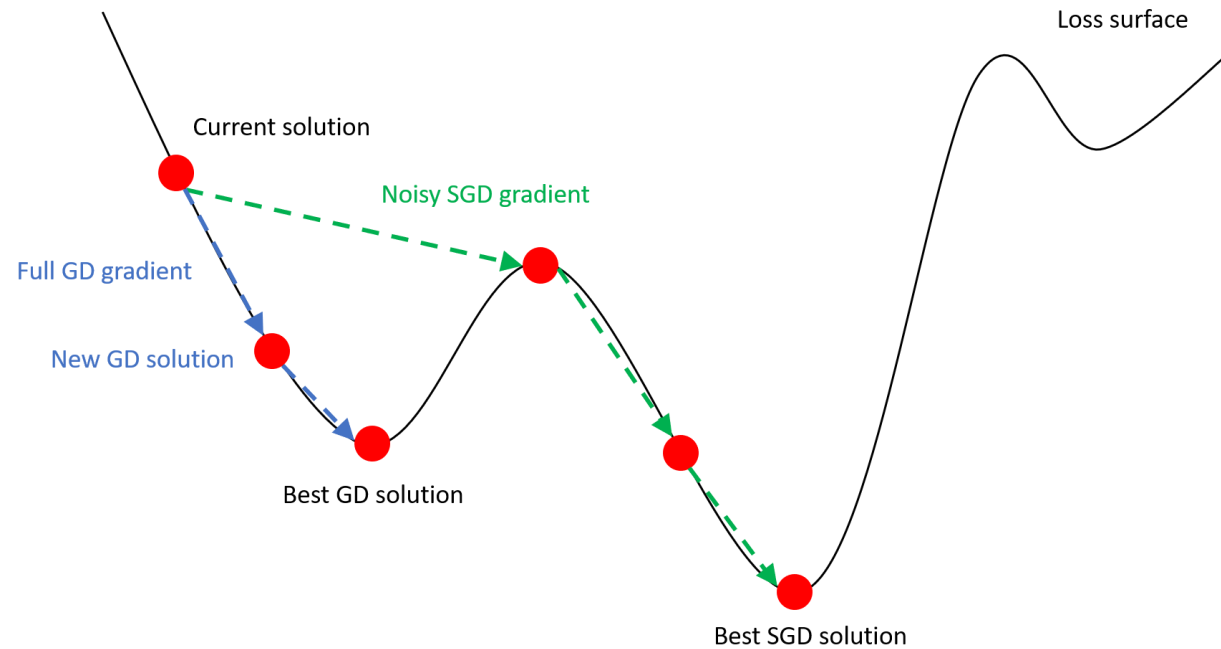
$$\omega \leftarrow \omega - \alpha \frac{\partial \ell^{(i)}}{\partial \omega}$$



- SGD takes steps in a noisy direction, but moves downhill on average.
- Mathematical justification: if you sample a training example at random, the stochastic gradient is an unbiased estimate of the batch gradient:

$$\mathbb{E} \left[\frac{\partial \ell^{(i)}}{\partial \omega} \right] = \frac{1}{m} \sum_{i=1}^m \frac{\partial \ell^{(i)}}{\partial \omega} = \frac{\partial}{\partial \omega} \left[\frac{1}{m} \sum_{i=1}^m \ell^{(i)} \right] = \frac{\partial \mathcal{E}}{\partial \omega}$$

SGD is Sometimes Better



- No guarantee that this is what is going to always happen.
- But the noisy SGD gradients can help occasionally escaping local optima

Mini-batch Gradient Descent

- Potential problem of SGD: gradient estimates can be very noisy
- Compromise approach: compute the gradients on a medium-sized set of training examples s ($< m$), called a **mini-batch**.

$$\mathcal{E}(\omega) = \frac{1}{s} \sum_{i=1}^s \ell(\hat{y}_i, y_i) = \frac{1}{s} \sum_{i=1}^s \ell(h_{\omega}(x_i), y_i) = \frac{1}{s} \sum_{i=1}^s \ell^{(i)}$$

$$\omega \leftarrow \omega - \alpha \nabla_{\omega} \mathcal{E}$$

- Stochastic gradients computed on larger mini-batches have smaller variance:


$$\text{var} \left[\frac{1}{s} \sum_{i=1}^s \frac{\partial \ell^{(i)}}{\partial \omega} \right] = \frac{1}{s^2} \text{var} \left[\sum_{i=1}^s \frac{\partial \ell^{(i)}}{\partial \omega} \right] = \frac{1}{s} \text{var} \left[\frac{\partial \ell^{(i)}}{\partial \omega} \right]$$

- The mini-batch size s is a hyper-parameter that needs to be set.

Implementation with TensorFlow

Batch Gradient Descent with TensorFlow

- We will explore the python codes of these three gradient descent algorithms with a logistic regression model.



Batch Gradient Descent with TensorFlow

```
LR = 0.04
n_iter = 60000
n_prt = 250

x = tf.placeholder(tf.float32, [m, 3])
y = tf.placeholder(tf.float32, [m, 1])

w = tf.Variable([[0],[0],[0]], dtype = tf.float32)

y_pred = tf.matmul(x,w)
loss = tf.nn.sigmoid_cross_entropy_with_logits(logits=y_pred, labels=y)
loss = tf.reduce_mean(loss)

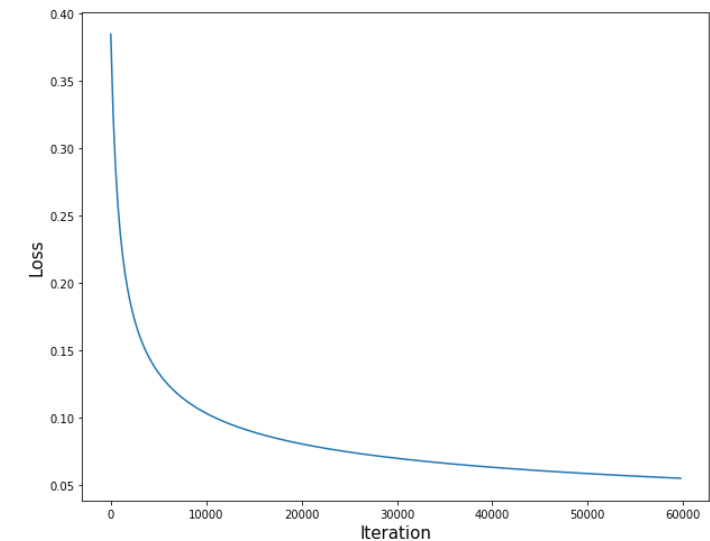
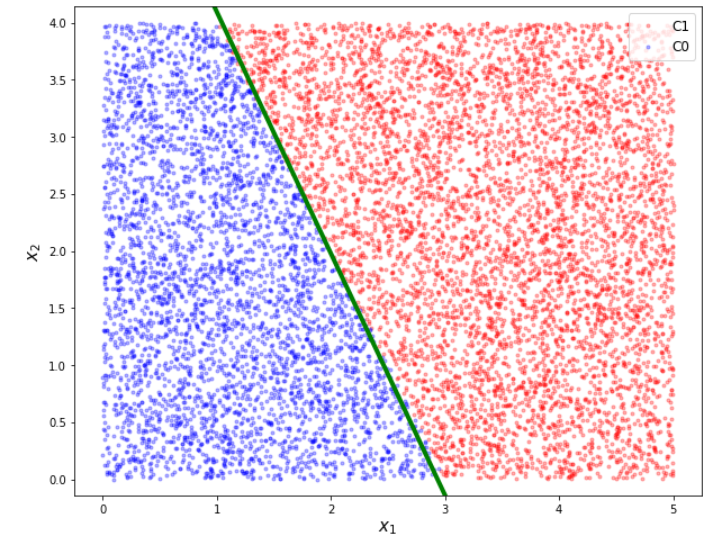
optm = tf.train.GradientDescentOptimizer(LR).minimize(loss)
init = tf.global_variables_initializer()

start_time = time.time()

loss_record = []
with tf.Session() as sess:
    sess.run(init)
    for epoch in range(n_iter):
        sess.run(optm, feed_dict = {x: train_X, y: train_y})

        if (epoch + 1) % n_prt == 0:
            loss_record.append(sess.run(loss, feed_dict = {x: train_X, y: train_y}))

    w_hat = sess.run(w)
```



Stochastic Gradient Descent (SGD) with TensorFlow

```
LR = 0.04
n_iter = 60000
n_prt = 250

x = tf.placeholder(tf.float32, [1, 3])
y = tf.placeholder(tf.float32, [1, 1])

w = tf.Variable(tf.random_normal([3,1]), dtype = tf.float32)

y_pred = tf.matmul(x,w)
loss = tf.nn.sigmoid_cross_entropy_with_logits(logits=y_pred, labels=y)
loss = tf.reduce_mean(loss)

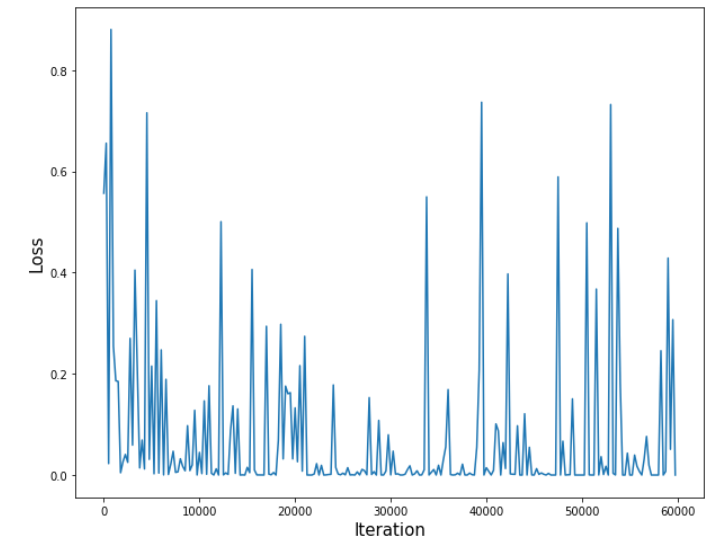
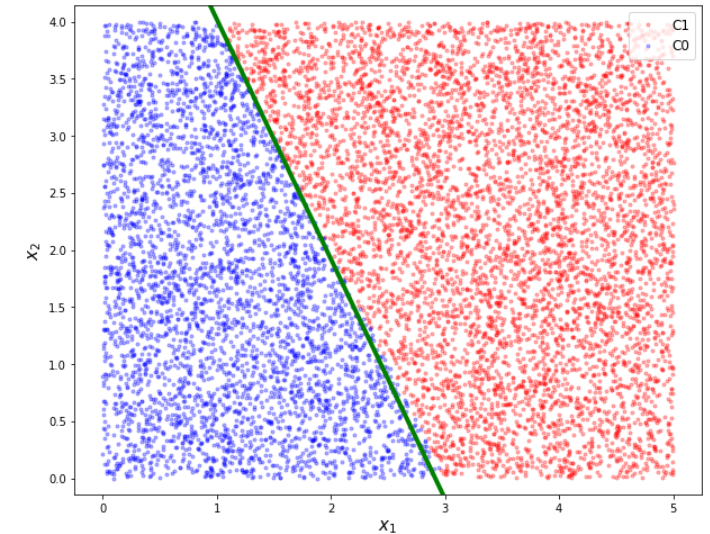
optm = tf.train.GradientDescentOptimizer(LR).minimize(loss)
init = tf.global_variables_initializer()

start_time = time.time()

loss_record = []
with tf.Session() as sess:
    sess.run(init)
    for epoch in range(n_iter):
        idx = np.random.choice(m, 1)
        batch_X = train_X[idx,:]
        batch_y = train_y[idx]
        sess.run(optm, feed_dict = {x: batch_X, y: batch_y})

        if (epoch + 1) % n_prt == 0:
            loss_record.append(sess.run(loss, feed_dict = {x: batch_X, y: batch_y}))

    w_hat = sess.run(w)
```



Mini-batch Gradient Descent with TensorFlow

```
LR = 0.04
n_iter = 60000
n_batch = 50
n_prt = 250

x = tf.placeholder(tf.float32, [n_batch, 3])
y = tf.placeholder(tf.float32, [n_batch, 1])

w = tf.Variable(tf.random_normal([3,1]), dtype = tf.float32)

y_pred = tf.matmul(x,w)
loss = tf.nn.sigmoid_cross_entropy_with_logits(logits=y_pred, labels=y)
loss = tf.reduce_mean(loss)

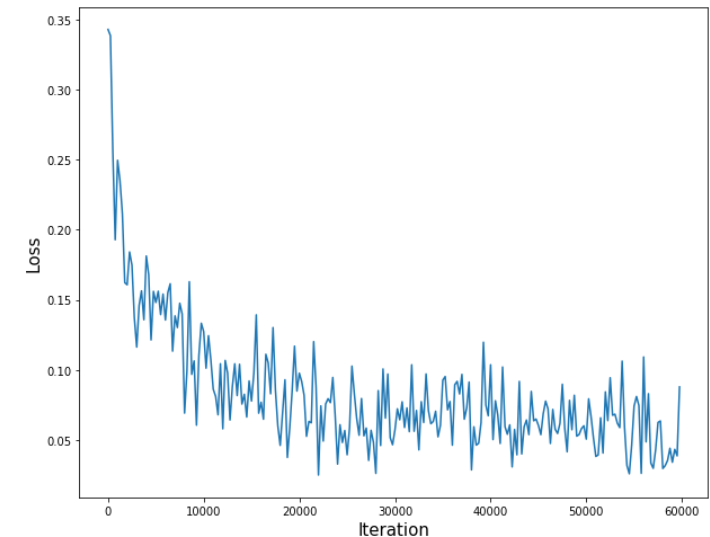
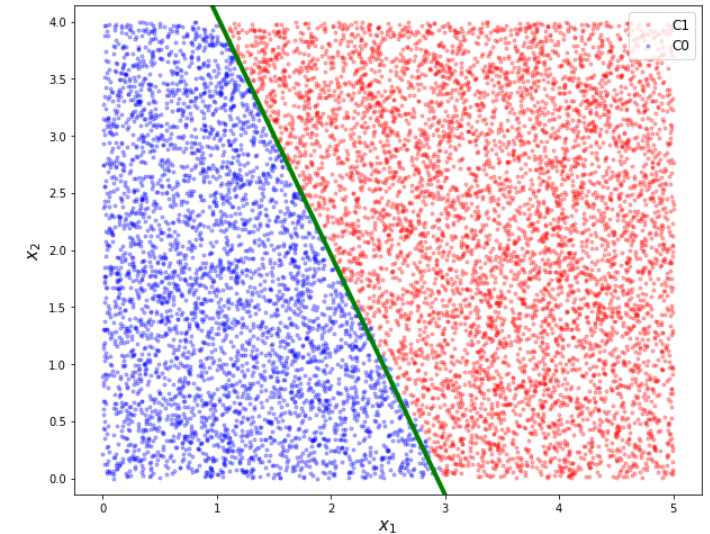
optm = tf.train.GradientDescentOptimizer(LR).minimize(loss)
init = tf.global_variables_initializer()

start_time = time.time()

loss_record = []
with tf.Session() as sess:
    sess.run(init)
    for epoch in range(n_iter):
        idx = np.random.choice(m, size = n_batch)
        batch_X = train_X[idx,:]
        batch_y = train_y[idx]
        sess.run(optm, feed_dict = {x: batch_X, y: batch_y})

        if (epoch + 1) % n_prt == 0:
            loss_record.append(sess.run(loss, feed_dict = {x: batch_X, y: batch_y}))

    w_hat = sess.run(w)
```



Limitation of the Gradient Descent

Setting the Learning Rate


- How can we set the learning rate?

$$\omega_{k+1} = \omega_k - \underset{\uparrow}{\alpha} \nabla f(\omega_k)$$

- Small learning rate converges slowly and gets stuck in false local minima
- Large learning rates overshoot, become unstable and diverge
- Idea 1
 - Try lots of different learning rates and see what works “just right”
- Idea 2
 - Do something smarter! Design an adaptive learning rate that “adapts” to the landscape
 - Temporal and spatial

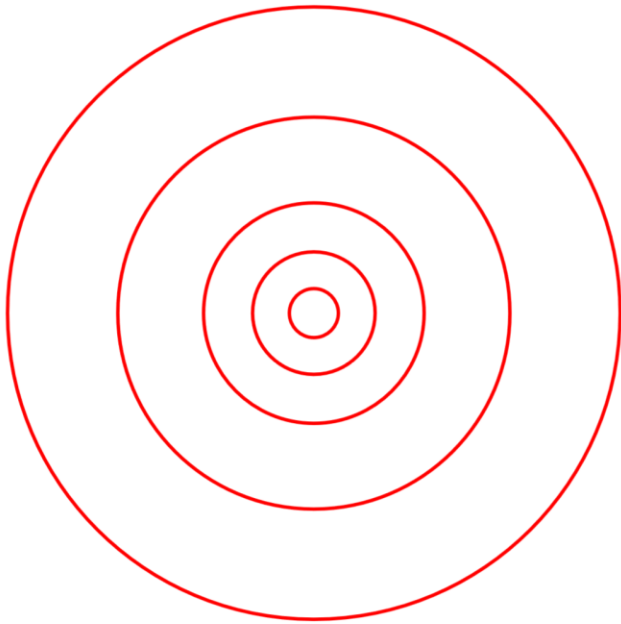
SGD Learning Rate (= Step Size)

- The gradient descent method makes a strong assumption about the magnitude of the “local curvature” to fix the step size, and about its isotropy so that the same step size makes sense in all directions.

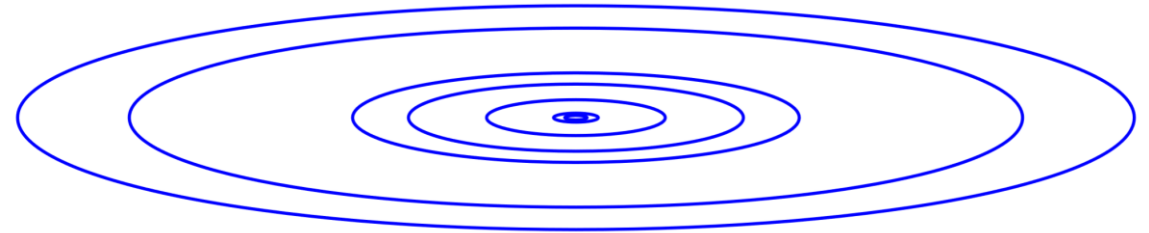


SGD Learning Rate: Spatial

- We assign the same learning rate to all features



Nice (all features are equally important)

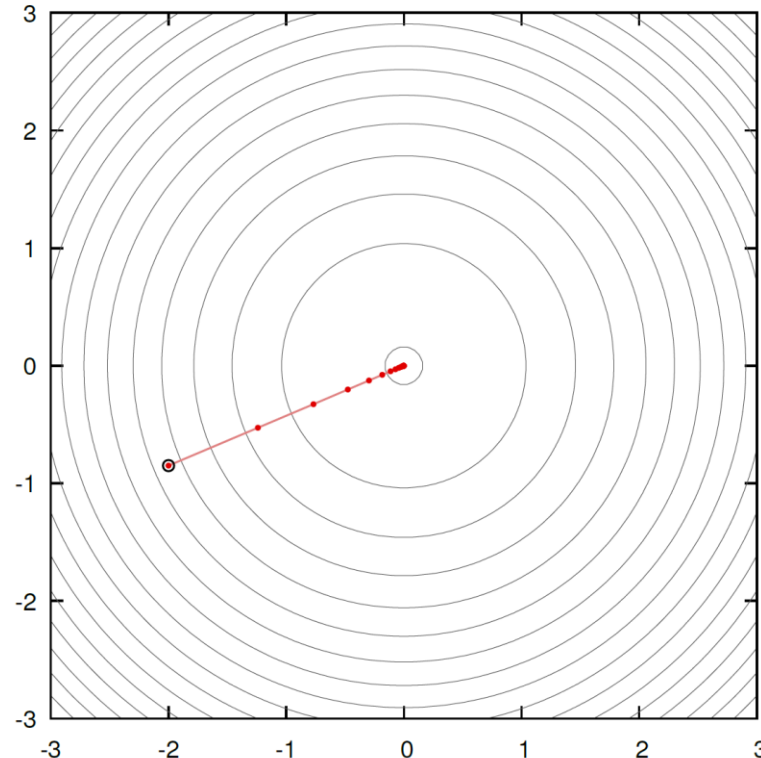


Harder !

SGD Learning Rate: Spatial

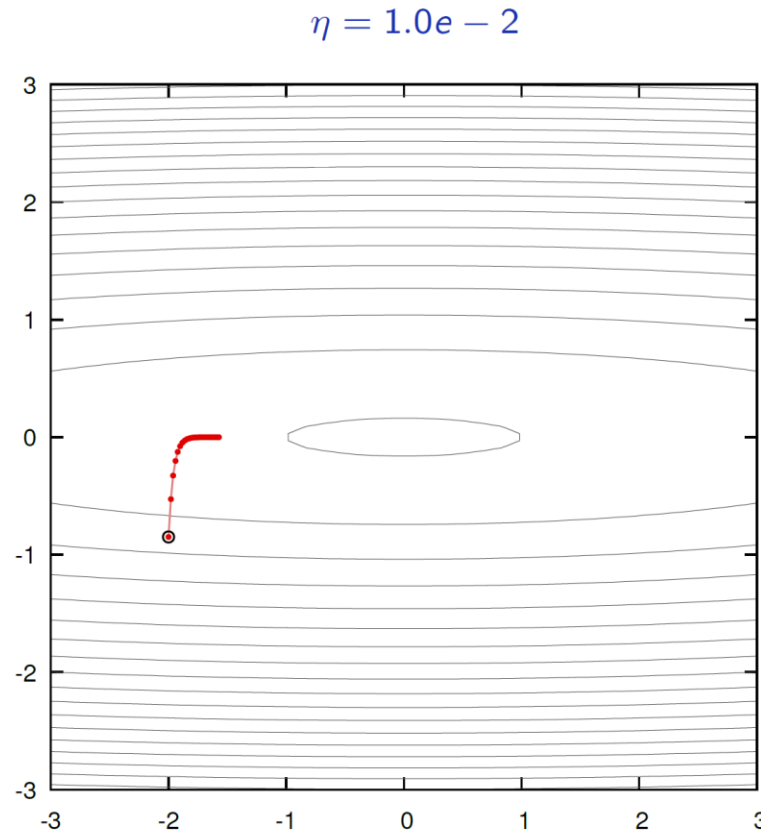
- The gradient descent method makes a strong assumption about the magnitude of the “local curvature” to fix the step size, and about its isotropy so that the same step size makes sense in all directions.
- Nice (all features are equally important)

$$\eta = 1.0e - 2$$



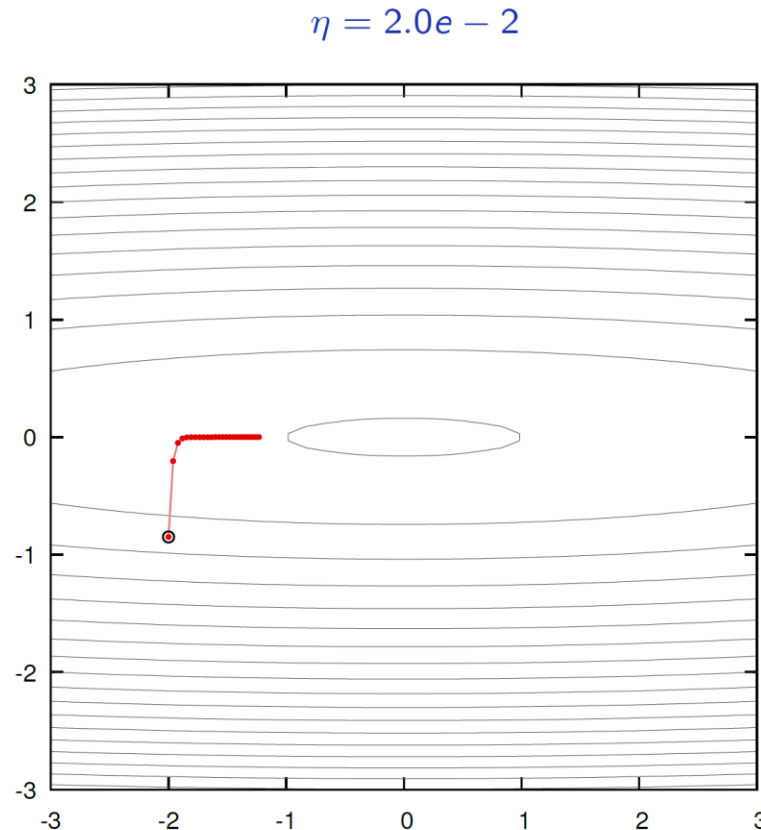
SGD Learning Rate: Spatial

- The gradient descent method makes a strong assumption about the magnitude of the “local curvature” to fix the step size, and about its isotropy so that the same step size makes sense in all directions.
- Harder !




SGD Learning Rate: Spatial

- The gradient descent method makes a strong assumption about the magnitude of the “local curvature” to fix the step size, and about its isotropy so that the same step size makes sense in all directions.
- Harder !




SGD Learning Rate: Spatial

- The gradient descent method makes a strong assumption about the magnitude of the “local curvature” to fix the step size, and about its isotropy so that the same step size makes sense in all directions.
- Harder !




SGD Learning Rate: Spatial

- The gradient descent method makes a strong assumption about the magnitude of the “local curvature” to fix the step size, and about its isotropy so that the same step size makes sense in all directions.
- Harder !



SGD Learning Rate: Spatial

- The gradient descent method makes a strong assumption about the magnitude of the “local curvature” to fix the step size, and about its isotropy so that the same step size makes sense in all directions.
- Harder !



SGD Learning Rate: Temporal

- Typical strategy:
 - Use a large learning rate early in training so you can get close to the optimum
 - Gradually **decay** the learning rate to reduce the fluctuations

Adaptive Learning Rate Methods


- SDG

$$\theta_{t+1,i} = \theta_{t,i} - \eta \cdot g_{t,i}.$$

- Adagrad

$$\theta_{t+1,i} = \theta_{t,i} - \frac{\eta}{\sqrt{G_{t,ii} + \epsilon}} \cdot g_{t,i}.$$

- $G_{t,ii}$ sum of the squares of the gradients
 - Perform smaller updates (i.e. low learning rates) for parameters associated with frequent occurring features, and
 - Perform larger updates (i.e. high learning rates) for parameters associated with infrequent features
- Deep-learning generally relies on a smarter use of the gradient, using statistics over its past values to make a “smarter step” with the current one.



Adaptive Learning Rate Methods

- Momentum
- Adagrad
- Adadelata
- Adam
- RMSProp



`tf.train.MomentumOptimizer`



`tf.train.AdagradOptimizer`



`tf.train.AdadeltaOptimizer`



`tf.train.AdamOptimizer`



`tf.train.RMSPropOptimizer`

Qian et al. "On the momentum term in gradient descent learning algorithms." 1999.

Duchi et al. "Adaptive Subgradient Methods for Online Learning and Stochastic Optimization." 2011.

Zeiler et al. "ADADELTA: An Adaptive Learning Rate Method." 2012.

Kingma et al. "Adam: A Method for Stochastic Optimization." 2014.

- Additional detail: <http://runder.io/optimizing-gradient-descent/>