



Autoencoder

Prof. Seungchul Lee
Industrial AI Lab.

Unsupervised Learning

- Definition
 - Unsupervised learning refers to most attempts to extract information from a distribution that do not require human labor to annotate example
 - Main task is to find the ‘best’ representation of the data
- Dimension Reduction
 - Attempt to compress as much information as possible in a smaller representation
 - Preserve as much information as possible while obeying some constraint aimed at keeping the representation simpler
 - This modeling consists of finding “meaningful degrees of freedom” that describe the signal, and are of lesser dimension.

Autoencoders

- It is like 'deep learning version' of unsupervised learning
- Definition
 - An autoencoder is a neural network that is trained to attempt to copy its input to its output
 - The network consists of two parts: an encoder and a decoder that produce a reconstruction
- Encoder and Decoder
 - Encoder function : $z = f(x)$
 - Decoder function : $x = g(z)$
 - We learn to set $g(f(x)) = x$


Autoencoder

- Dimension reduction
- Recover the input data



Autoencoder

- Dimension reduction
- Recover the input data
 - Learns an encoding of the inputs so as to recover the original input from the encodings as well as possible




Original space

Latent space

Autoencoder

- Autoencoder combines an encoder f from the original space \mathcal{X} to a latent space \mathcal{F} , and a decoder g to map back to \mathcal{X} , such that $g \circ f$ is [close to] the identity on the data



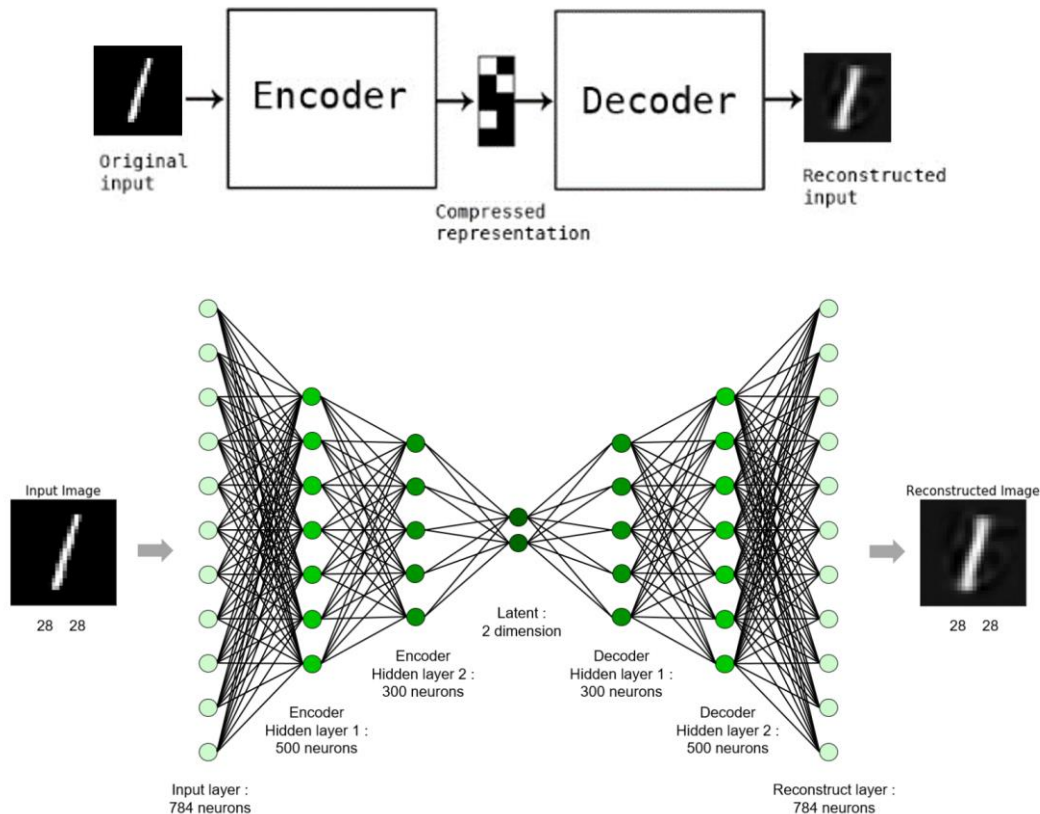
$$\mathbb{E} [\|X - g \circ f(X)\|^2] \approx 0$$

- A proper autoencoder has to capture a "good" parametrization of the signal, and in particular the statistical dependencies between the signal components.

Autoencoder with MNIST

Autoencoder with TensorFlow

- MNIST example
- Use only (1, 5, 6) digits to visualize in 2-D




$$\frac{1}{m} \sum_{i=1}^m (t_i - y_i)^2$$


Test or Evaluation

```
test_x, _ = test_batch_maker(1)  
x_reconst = sess.run(reconst, feed_dict = {x: test_x})
```

Input Image




Reconstructed Image



Distribution in Latent Space

- Make a projection of 784-dim image onto 2-dim latent space


```
test_x, test_y = test_batch_maker(500)
test_y = np.argmax(test_y, axis = 1)
test_latent = sess.run(latent, feed_dict = {x: test_x})
```




Autoencoder as Generative Model

Generative Capabilities


- We can assess the generative capabilities of the decoder g by introducing a [simple] density model q^Z over the latent space \mathcal{F} , sample there, and map the samples into the image space \mathcal{X} with g .



MNIST Example




Generated Fake Image



Latent Representation


- To get an intuition of the latent representation, we can pick two samples x and x' at random and interpolate samples along the line in the latent space



Latent Representation

- To get an intuition of the latent representation, we can pick two samples x and x' at random and interpolate samples along the line in the latent space

$$g((1 - \alpha)f(x) + \alpha f(x'))$$




Interpolation in High Dimension

Reasonable distance metric



Interpolation in high dimension




<https://www.cs.cmu.edu/~efros/courses/AP06/presentations/ThompsonDimensionalityReduction.pdf>

Interpolation in Manifold

Reasonable distance metric





Interpolation in manifold



<https://www.cs.cmu.edu/~efros/courses/AP06/presentations/ThompsonDimensionalityReduction.pdf>

MNIST Example: Walk in the Latent Space



Generative Models

- It generates something that makes sense.
- These results are unsatisfying, because the density model used on the latent space \mathcal{F} is too simple and inadequate.
- Building a “good” model amounts to our original problem of modeling an empirical distribution, although it may now be in a lower dimension space.
- This is a motivation to VAE or GAN.