




Fully Convolutional Network (FCN)

Prof. Seungchul Lee
Industrial AI Lab.

Deep Learning for Computer Vision: Review


Foundations

- Why computer vision?
- Representing images
- Convolutions for feature extraction



CNNs

- CNN architecture
- Application to classification: ImageNet




Applications

- Segmentation, object detection, image captioning
- Visualization




Segmentation

- Segmentation task is different from classification task because it requires predicting a class for each pixel of the input image, instead of only 1 class for the whole input.
- Segment images into regions with different semantic categories. These semantic regions label and predict objects at the pixel level



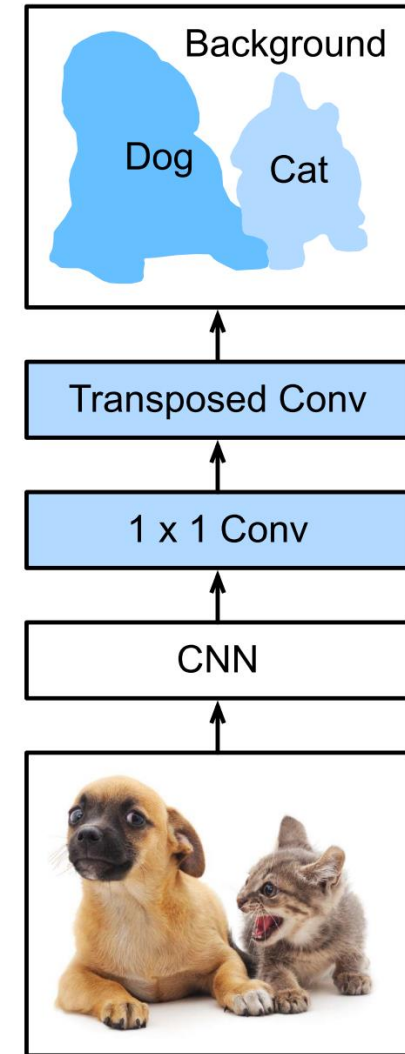
Segmentation

- Segmentation task is different from classification task because it requires predicting a class for each pixel of the input image, instead of only 1 class for the whole input.
- Segment images into regions with different semantic categories. These semantic regions label and predict objects at the pixel level
- Classification needs to understand what is in the input (namely, the context).
- However, in order to predict what is in the input for each pixel, segmentation needs to recover not only what is in the input, but also **where**.





Semantic Segmentation: FCNs

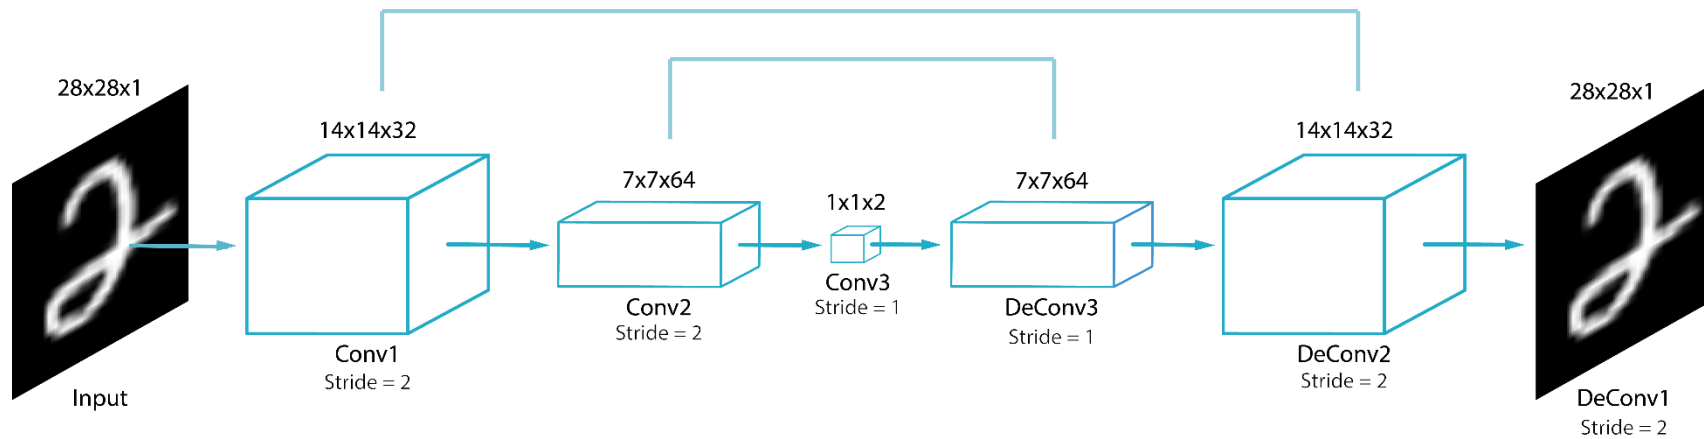

- FCN uses a convolutional neural network to transform image pixels to pixel categories.
- Network designed with all convolutional layers, with down-sampling and up-sampling operations
- Given a position on the spatial dimension, the output of the channel dimension will be a category prediction of the pixel corresponding to the location.



From CAE to FCN




From CAE to FCN



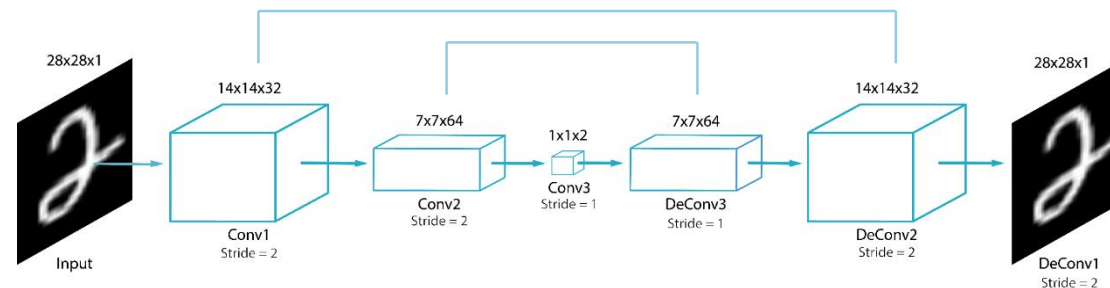
Skip Connection

- A skip connection is a connection that bypasses at least one layer.
- Here, it is often used to transfer local information by summing feature maps from the downsampling path with feature maps from the upsampling path.
 - Merging features from various resolution levels helps combining context information with spatial information.



Fully Convolutional Networks (FCNs)

- To obtain a segmentation map (output), segmentation networks usually have 2 parts
 - Downsampling path: capture semantic/contextual information
 - Upsampling path: recover spatial information
- The downsampling path is used to extract and interpret the context (what), while the upsampling path is used to enable precise localization (where).
- Furthermore, to fully recover the fine-grained spatial information lost in the pooling or downsampling layers, we often use **skip connections**.
- Network can work regardless of the original image size, without requiring any fixed number of units at any stage.



Segmented (Labeled) Images

input




output




output



FCN Architecture



FCN Architecture



```

vgg16_weights = model.get_weights()

weights = {
    'conv1_1' : tf.constant(vgg16_weights[0]),
    'conv1_2' : tf.constant(vgg16_weights[2]),

    'conv2_1' : tf.constant(vgg16_weights[4]),
    'conv2_2' : tf.constant(vgg16_weights[6]),

    'conv3_1' : tf.constant(vgg16_weights[8]),
    'conv3_2' : tf.constant(vgg16_weights[10]),
    'conv3_3' : tf.constant(vgg16_weights[12]),

    'conv4_1' : tf.constant(vgg16_weights[14]),
    'conv4_2' : tf.constant(vgg16_weights[16]),
    'conv4_3' : tf.constant(vgg16_weights[18]),

    'conv5_1' : tf.constant(vgg16_weights[20]),
    'conv5_2' : tf.constant(vgg16_weights[22]),
    'conv5_3' : tf.constant(vgg16_weights[24]),
}

```

```

biases = {
    'conv1_1' : tf.constant(vgg16_weights[1]),
    'conv1_2' : tf.constant(vgg16_weights[3]),

    'conv2_1' : tf.constant(vgg16_weights[5]),
    'conv2_2' : tf.constant(vgg16_weights[7]),


    'conv3_1' : tf.constant(vgg16_weights[9]),
    'conv3_2' : tf.constant(vgg16_weights[11]),
    'conv3_3' : tf.constant(vgg16_weights[13]),

    'conv4_1' : tf.constant(vgg16_weights[15]),
    'conv4_2' : tf.constant(vgg16_weights[17]),
    'conv4_3' : tf.constant(vgg16_weights[19]),

    'conv5_1' : tf.constant(vgg16_weights[21]),
    'conv5_2' : tf.constant(vgg16_weights[23]),
    'conv5_3' : tf.constant(vgg16_weights[25]),
}

```

FCN Architecture




```
def fcn(x, weights, biases):
    # First convolution layers
    conv1_1 = tf.nn.conv2d(x,
                           weights['conv1_1'],
                           strides = [1, 1, 1, 1],
                           padding = 'SAME')
    conv1_1 = tf.nn.relu(tf.add(conv1_1, biases['conv1_1']))
    conv1_2 = tf.nn.conv2d(conv1_1,
                           weights['conv1_2'],
                           strides = [1, 1, 1, 1],
                           padding = 'SAME')
    conv1_2 = tf.nn.relu(tf.add(conv1_2, biases['conv1_2']))
    maxp1 = tf.nn.max_pool(conv1_2,
                           ksize = [1, 2, 2, 1],
                           strides = [1, 2, 2, 1],
                           padding = 'VALID')

    # Second convolution layers
    conv2_1 = tf.nn.conv2d(maxp1,
                           weights['conv2_1'],
                           strides = [1, 1, 1, 1],
                           padding = 'SAME')
    conv2_1 = tf.nn.relu(tf.add(conv2_1, biases['conv2_1']))
    conv2_2 = tf.nn.conv2d(conv2_1,
                           weights['conv2_2'],
                           strides = [1, 1, 1, 1],
                           padding = 'SAME')
    conv2_2 = tf.nn.relu(tf.add(conv2_2, biases['conv2_2']))
    maxp2 = tf.nn.max_pool(conv2_2,
                           ksize = [1, 2, 2, 1],
                           strides = [1, 2, 2, 1],
                           padding = 'VALID')
```

Fixed

Trained

FCN Architecture



```
# sixth convolution layer
conv6 = tf.layers.conv2d(maxp5,
                          filters = 4096,
                          kernel_size = 7,
                          padding = 'SAME',
                          activation = tf.nn.relu)

# 1x1 convolution layers
fc4 = tf.layers.conv2d(conv6,
                       filters = 4096,
                       kernel_size = 1,
                       padding = 'SAME',
                       activation = tf.nn.relu)

fc3 = tf.layers.conv2d(fc4,
                       filters = 2,
                       kernel_size = 1,
                       padding = 'SAME')
```

```
# Upsampling layers
fc2 = tf.layers.conv2d_transpose(fc3,
                                 filters = 512,
                                 kernel_size = 4,
                                 strides = (2, 2),
                                 padding = 'SAME')

fc1 = tf.layers.conv2d_transpose(fc2 + maxp4,
                                 filters = 256,
                                 kernel_size = 4,
                                 strides = (2, 2),
                                 padding = 'SAME')

output = tf.layers.conv2d_transpose(fc1 + maxp3,
                                     filters = 2,
                                     kernel_size = 16,
                                     strides = (8, 8),
                                     padding = 'SAME')
```


Segmentation Result

