



Unsupervised Learning: Dimension Reduction

Prof. Seungchul Lee
Industrial AI Lab.

Dimension Reduction


- Motivation:
 - Can we describe high-dimensional data in a “simpler” way?
- Dimension reduction without losing too much information
- Find a low-dimensional, yet useful representation of the data

Dimension Reduction



- Why dimensionality reduction?
 - insights into the low-dimensional structures in the data (visualization)
 - Fewer dimensions \Rightarrow Less chances of overfitting \Rightarrow Better generalization
 - **Speeding** up learning algorithms
 - Most algorithms scale badly with increasing data dimensionality
 - **Less storage** requirements (data compression)
 - Note: Dimensionality reduction is **different from feature selection**
 - ... although the goals are kind of the same
 - Dimensionality reduction is more like “**feature extraction**”
 - Constructing a small set of new features from the original features

Highly Correlated Data

- How?
 - idea: highly correlated data contains redundant features




low redundancy



high redundancy

Principal Component Analysis (PCA)


- Each example x has 2 features $\{x_1, x_2\}$
- Consider ignoring the feature x_2 for each example
- Each 2-dimensional example x now becomes 1-dimensional $x = \{x_1\}$
- Are we losing much information by throwing away x_2 ?



- **No.** Most of the data spread is along x_1 (very little variance along x_2)

Principal Component Analysis (PCA)


- Each example x has 2 features $\{x_1, x_2\}$
- Consider ignoring the feature x_2 for each example
- Each 2-dimensional example x now becomes 1-dimensional $x = \{x_1\}$
- Are we losing much information by throwing away x_2 ?



- Yes, the data has substantial variance along both features (i.e., both axes)

Principal Component Analysis (PCA)


- Now consider a change of axes
- Each example x has 2 features $\{u_1, u_2\}$
- Consider ignoring the feature u_2 for each example
- Each 2-dimensional example x now become 1-dimensional $x = \{u_1\}$




- **No.** Most of the data spread is along u_1 (very little variance along u_2)

Principal Component Analysis (PCA)

- Data \rightarrow projection onto unit vector \hat{u}_1
 - PCA is used when we want projections capturing **maximum variance directions**
 - Principal Components (PC): directions u_1 of maximum variability in the data
 - Roughly speaking, PCA does a change of axes that can represent the data in a succinct manner



Principal Component Analysis (PCA)



- HOW?
 1. Maximize variance (most separable)
 2. Minimize the sum-of-squares (minimum squared error)

PCA Algorithm: Pre-processing

- Given data

$$\mathbf{x}^{(i)} = \begin{bmatrix} x_1^{(i)} \\ \vdots \\ x_n^{(i)} \end{bmatrix}, \quad \mathbf{X} = \begin{bmatrix} \dots & (\mathbf{x}^{(1)})^T & \dots \\ \dots & (\mathbf{x}^{(2)})^T & \dots \\ & \vdots & \\ \dots & (\mathbf{x}^{(m)})^T & \dots \end{bmatrix}$$

- Shifting (zero mean) and rescaling (unit variance)

1. Shift to zero mean

$$\mu = \frac{1}{m} \sum_{i=1}^m \mathbf{x}^{(i)}$$
$$\mathbf{x}^{(i)} \leftarrow \mathbf{x}^{(i)} - \mu \quad (\text{zero mean})$$

2. [optional] Rescaling (unit variance)

$$\sigma_j^2 = \frac{1}{m-1} \sum_{i=1}^m m \left(x_j^{(i)} \right)^2$$
$$x_j^{(i)} \leftarrow \frac{x_j^{(i)}}{\sigma_j}$$

PCA Algorithm: Maximize Variance

- Find unit vector u such that **maximizes variance of projections**
 - Note: $m \approx m - 1$ for big data

$$\begin{aligned}\text{variance of projected data} &= \frac{1}{m} \sum_{i=1}^m (u^T x^{(i)})^2 = \frac{1}{m} \sum_{i=1}^m (x^{(i)T} u)^2 \\ &= \frac{1}{m} \sum_{i=1}^m (x^{(i)T} u)^T (x^{(i)T} u) = \frac{1}{m} \sum_{i=1}^m u^T x^{(i)} x^{(i)T} u \\ &= u^T \left(\frac{1}{m} \sum_{i=1}^m x^{(i)} x^{(i)T} \right) u \\ &= u^T S u \quad \left(S = \frac{1}{m} X^T X : \text{sample covariance matrix} \right)\end{aligned}$$

Maximize Variance

- In an optimization form

$$\begin{array}{ll}\text{maximize} & u^T S u \\ \text{subject to} & u^T u = 1\end{array}$$


$$u^T S u = u^T \lambda u = \lambda u^T u = \lambda \quad (\text{Eigen analysis : } S u = \lambda u)$$

\implies pick the largest eigenvalue λ_1 of covariance matrix S

$\implies u = u_1$ is the λ_1 's corresponding eigenvector

$\implies u_1$ is the first principal component (direction of highest variance in the data)

Minimize the Sum-of-Squared Error



$$\begin{aligned}\|e^{(i)}\|^2 &= \|x^{(i)}\|^2 - (x^{(i)T} u)^2 \\ &= \|x^{(i)}\|^2 - (x^{(i)T} u)^T (x^{(i)T} u) \\ &= \|x^{(i)}\|^2 - u^T x^{(i)} x^{(i)T} u\end{aligned}$$

$$\begin{aligned}\frac{1}{m} \sum_{i=1}^m \|e^{(i)}\|^2 &= \frac{1}{m} \sum_{i=1}^m \|x^{(i)}\|^2 - \frac{1}{m} \sum_{i=1}^m u^T x^{(i)} x^{(i)T} u \\ &= \frac{1}{m} \sum_{i=1}^m \|x^{(i)}\|^2 - u^T \left(\frac{1}{m} \sum_{i=1}^m x^{(i)} x^{(i)T} \right) u\end{aligned}$$

Minimize the Sum-of-Squared Error

- In an optimization form

$$\text{minimize } \underbrace{\frac{1}{m} \sum_{i=1}^m \|x^{(i)}\|^2}_{\text{constant given } x_i} - \underbrace{u^T \left(\frac{1}{m} \sum_{i=1}^m x^{(i)} x^{(i)T} \right) u}_{\text{maximize}}$$

$$\implies \text{maximize } u^T \left(\frac{1}{m} \sum_{i=1}^m x^{(i)} x^{(i)T} \right) u = \max u^T S u$$

$$\therefore \text{minimize } error^2 = \text{maximize } variance$$

Dimension Reduction Method ($n \rightarrow k$)

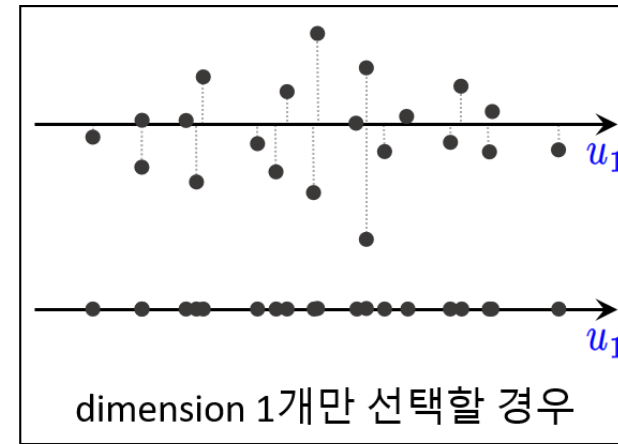


1. Choose top k (orthonormal) eigenvectors, $U = [u_1, u_2, \dots, u_k]$
2. Project x_i onto span $\{u_1, u_2, \dots, u_k\}$

$$z^{(i)} = \begin{bmatrix} u_1^T x^{(i)} \\ u_2^T x^{(i)} \\ \vdots \\ u_k^T x^{(i)} \end{bmatrix} \quad \text{or} \quad z = U^T x$$


- $x^{(i)} \rightarrow$ projection onto unit vector $u \Rightarrow u^T x^{(i)} =$ distance from the origin along u

Principal Component Analysis

- Data \rightarrow projection onto unit vector u






Pictorial Summary of PCA



PCA Visualization

- Eigenvectors
 - Given basis $\{\hat{x}_1, \hat{x}_2\}$ to transformed basis $\{\hat{u}_1, \hat{u}_2\}$


$$[\hat{u}_1 \ \hat{u}_2] = [\hat{x}_1 \ \hat{x}_2] \begin{bmatrix} c_1 & c_3 \\ c_2 & c_4 \end{bmatrix}$$




- Eigenvalues
 - λ_1, λ_2 indicates variance along the eigenvectors, respectively.
 - The larger eigenvalue is, the more dominant feature (eigenvector) is.

Linear Regression vs. PCA

Linear Regression




PCA



Python Codes

```
# data generation
m = 5000
mu = np.array([0, 0])
sigma = np.array([[3, 1.5],
                  [1.5, 1]])

X = np.random.multivariate_normal(mu, sigma, m)
X = np.asmatrix(X)
```



Python Codes

```
S = 1/(m-1)*X.T*X
S = np.asmatrix(S)

D, U = np.linalg.eig(S)

idx = np.argsort(-D)
D = D[idx]
U = U[:,idx]


print(D, '\n')
print(U)
```

```
[3.90595228 0.19884608]
```

```
[[ 0.88112544 -0.4728826 ]
 [ 0.4728826  0.88112544]]
```


```
h = U[1,0]/U[0,0]
xp = np.arange(-6, 6, 0.1)
yp = h*xp
```

maximize $u^T S u$
subject to $u^T u = 1$




```
from sklearn.decomposition import PCA  
  
pca = PCA(n_components = 1)  
pca.fit(X)
```

```
u = pca.transform(X)  
  
plt.figure(figsize = (10, 8))  
plt.hist(u, 51)  
plt.show()
```



PCA Example


- Multiple video camera records of spring and mass system
- Optimal data representation
 - Find the most informative point of view




- source:
 - https://www.cs.princeton.edu/picasso/mats/PCA-Tutorial-Intuition_jp.pdf

Multivariate Time Series

- System order can be inferred from
 - Laws of physics or
 - Data



Measured observations




Multivariate Time Series

$$x^{(i)} = \begin{bmatrix} x \text{ in camera 1} \\ y \text{ in camera 1} \\ x \text{ in camera 2} \\ y \text{ in camera 2} \\ x \text{ in camera 3} \\ y \text{ in camera 3} \end{bmatrix},$$


$$X_{m \times 6} = \begin{bmatrix} \dots & (x^{(1)})^T & \dots \\ \dots & (x^{(2)})^T & \dots \\ & \vdots & \\ & \vdots & \\ \dots & (x^{(m)})^T & \dots \end{bmatrix}$$

$$\begin{aligned} &\text{maximize} && u^T S u \\ &\text{subject to} && u^T u = 1 \end{aligned}$$




Eigenvalues

```
plt.figure(figsize = (10,8))  
plt.stem(np.sqrt(D))  
plt.grid(alpha = 0.3)  
plt.show()
```



Projection onto Principal Components

```
# relative magnitudes of the principal components  
  
Z = X*U  
xp = np.arange(0, m)/24    # 24 frame rate
```



PCA Example

