




Artificial Neural Networks


Prof. Seungchul Lee
Industrial AI Lab.

Perceptron: Example




$$\begin{aligned}\hat{y} &= g(\omega_0 + X^T \omega) \\ &= g\left(1 + \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}^T \begin{bmatrix} 3 \\ -2 \end{bmatrix}\right) \\ &= g(1 + 3x_1 - 2x_2)\end{aligned}$$


Perceptron: Example



$$\hat{y} = g(1 + 3x_1 - 2x_2)$$




Perceptron: Example




$$\hat{y} = g(1 + 3 \times (-1) - 2 \times 2) = g(-6) = -1$$


$$\hat{y} = g(1 + 3x_1 - 2x_2)$$




Perceptron: Example



$$\hat{y} = g(1 + 3x_1 - 2x_2)$$



Perceptron: Forward Propagation




$$\hat{y} = g(\omega_0 + X^T \omega)$$

$$= g\left(\omega_0 + \begin{bmatrix} x_1 \\ \vdots \\ x_m \end{bmatrix}^T \begin{bmatrix} \omega_1 \\ \vdots \\ \omega_m \end{bmatrix}\right)$$

From Perceptron to MLP

Artificial Neural Networks: Perceptron


- Perceptron for $h(\theta)$ or $h(\omega)$
 - Neurons compute the weighted sum of their inputs
 - A neuron is activated or fired when the sum a is positive



- A step function is not differentiable
- One neuron is often not enough
 - One hyperplane

$$a = \omega_0 + \omega_1 x_1 + \omega_2 x_2$$

$$\hat{y} = g(a) = \begin{cases} 1 & a > 0 \\ 0 & \text{otherwise} \end{cases}$$






Here, a step function is illustrated instead of a sign function

XOR Problem

- Minsky-Papert Controversy on XOR
 - Not linearly separable
 - Limitation of perceptron


x_1	x_2	$x_1 \text{ XOR } x_2$
0	0	0
0	1	1
1	0	1
1	1	0



- Single neuron = **one linear classification boundary**





Artificial Neural Networks: MLP

- Multi-layer Perceptron (MLP) = Artificial Neural Networks (ANN)
 - Multi neurons = multiple linear classification boundaries




Artificial Neural Networks: Activation Function

- Differentiable nonlinear activation function




Common Activation Functions

Sigmoid Function




$$g(z) = \frac{1}{1 + e^{-z}}$$

$$g'(z) = g(z)(1 - g(z))$$

 `tf.nn.sigmoid(z)`

Hyperbolic Tangent



$$g(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$


$$g'(z) = 1 - g(z)^2$$

 `tf.nn.tanh(z)`

Discuss later



Rectified Linear Unit (ReLU)




$$g(z) = \max(0, z)$$

$$g'(z) = \begin{cases} 1, & z > 0 \\ 0, & \text{otherwise} \end{cases}$$

 `tf.nn.relu(z)`


Artificial Neural Networks

- In a compact representation



Artificial Neural Networks


- A single layer is not enough to be able to represent complex relationship between input and output
⇒ perceptron with many layers and units



- Multi-layer perceptron
 - Features of features
 - Mapping of mappings

Another Perspective: ANN as Kernel Learning

Nonlinear Classification




SVM with a polynomial
Kernel visualization

Created by:
Udi Aharoni

Neuron


- We can represent this “neuron” as follows:

$$f(x) = \sigma(w \cdot x + b)$$



XOR Problem


- The main weakness of linear predictors is their lack of capacity.
- For classification, the populations have to be linearly separable.



“xor”

Nonlinear Mapping


- The XOR example can be solved by pre-processing the data to make the two populations linearly separable.



Nonlinear Mapping

- The XOR example can be solved by pre-processing the data to make the two populations linearly separable.


$$\phi : (x_u, x_v) \rightarrow (x_u, x_v, x_u x_v)$$



Nonlinear Mapping

- The XOR example can be solved by pre-processing the data to make the two populations linearly separable.

$$\phi : (x_u, x_v) \rightarrow (x_u, x_v, x_u x_v)$$




Kernel

- Often we want to capture nonlinear patterns in the data
 - nonlinear regression: input and output relationship may not be linear
 - nonlinear classification: classes may not be separable by a linear boundary
- Linear models (e.g. linear regression, linear SVM) are not just rich enough
 - by mapping data to higher dimensions where it exhibits linear patterns
 - apply the linear model in the new input feature space
 - mapping = changing the feature representation
- Kernels: make linear model work in nonlinear settings

Kernel + Neuron


- Nonlinear mapping + neuron

$$\phi : (x_u, x_v) \rightarrow (x_u, x_v, x_u x_v)$$



Neuron + Neuron


- Nonlinear mapping can be represented by another neurons



- Nonlinear Kernel
 - Nonlinear activation functions



Multi Layer Perceptron

- Nonlinear mapping can be represented by another neurons
- We can generalize an MLP




Summary

- Universal function approximator
- Universal function classifier
- Parameterized




Artificial Neural Networks

- Complex/Nonlinear universal function approximator
 - Linearly connected networks
 - Simple nonlinear neurons



Deep Artificial Neural Networks



- Complex/Nonlinear universal function approximator
 - Linearly connected networks
 - Simple nonlinear neurons



Looking at Parameters





Example: Linear Classifier

- Perceptron tries to separate the two classes of data by dividing them with a line




Example: Neural Networks

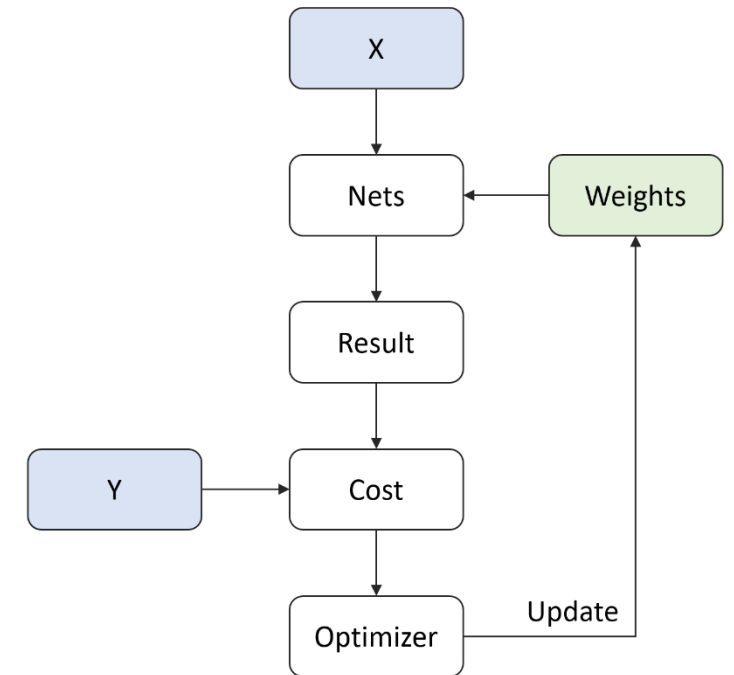

- The hidden layer learns a representation so that the data gets linearly separable



Logistic Regression in a Form of Neural Network




$$y = \sigma(\omega_0 + \omega_1 x_1 + \omega_2 x_2)$$




Logistic Regression in a Form of Neural Network

- Neural network convention


$$y = \sigma(\omega_0 + \omega_1 x_1 + \omega_2 x_2)$$




$$y = \sigma(b + \omega_1 x_1 + \omega_2 x_2)$$



Do not indicate bias units





```
n_input = 2  
n_output = 1
```



Logistic Regression in a Form of Neural Network


- One-hot encoding
 - One-hot encoding is a conventional practice for a multi-class classification

$$y^{(i)} \in \{1, 0\} \implies y^{(i)} \in \{[0, 1], [1, 0]\}$$




Nonlinearly Distributed Data


- Example to understand network's behavior
 - Include a hidden layer



Do not include bias units




```
n_input = 2  
n_hidden = 2  
n_output = 2
```




Multi Layers


- z space



Do not include bias units




```
n_input = 2  
n_hidden = 2  
n_output = 2
```




Multi Layers


- x space



Do not include bias units




```
n_input = 2  
n_hidden = 2  
n_output = 2
```




Nonlinearly Distributed Data


- More neurons in hidden layer



Do not include bias units




```
n_input = 2  
n_hidden = 4  
n_output = 2
```




Multi Layers


- Multiple linear classification boundaries



Do not include bias units



```
n_input = 2  
n_hidden = 4  
n_output = 2
```



(Artificial) Neural Networks: Training

Prof. Seungchul Lee
Industrial AI

Training Neural Networks: Optimization

- Learning or estimating weights and biases of multi-layer perceptron from training data
- 3 key components
 - objective function $f(\cdot)$
 - decision variable or unknown ω
 - constraints $g(\cdot)$
- In mathematical expression

$$\min_{\omega} f(\omega)$$

Training Neural Networks: Loss Function

- Measures error between target values and predictions

$$\min_{\omega} \sum_{i=1}^m \ell \left(h_{\omega} \left(x^{(i)} \right), y^{(i)} \right)$$

- Example

- Squared loss (for regression):

$$\frac{1}{m} \sum_{i=1}^m \left(h_{\omega} \left(x^{(i)} \right) - y^{(i)} \right)^2$$



- Cross entropy (for classification):

$$-\frac{1}{m} \sum_{i=1}^m y^{(i)} \log \left(h_{\omega} \left(x^{(i)} \right) \right) + \left(1 - y^{(i)} \right) \log \left(1 - h_{\omega} \left(x^{(i)} \right) \right)$$

Training Neural Networks: Gradient Descent



- Negative gradients points directly downhill of the cost function
- We can decrease the cost by moving in the direction of the negative gradient (α is a learning rate)

$$\omega \leftarrow \omega - \alpha \nabla_{\omega} \ell \left(h_{\omega} \left(x^{(i)} \right), y^{(i)} \right)$$



Gradients in ANN



- Learning weights and biases from data using gradient descent
- $\frac{\partial \ell}{\partial \omega}$: too many computations are required for all ω
- Structural constraint of NN:
 - Composition of functions
 - Chain rule
 - Dynamic programming



Dynamic Programming

Recursive Algorithm

- One of the central ideas of computer science
- Depends on solutions to smaller instances of the same problem (= sub-problem)
- Function to call itself (it is impossible in the real world)
- Factorial example
 - $n! = n \cdot (n - 1) \cdots 2 \cdot 1$



Dynamic Programming


- Dynamic Programming: general, powerful algorithm design technique
- Fibonacci numbers:

$$F_1 = F_2 = 1$$
$$F_n = F_{n-1} + F_{n-2}$$

Naïve Recursive Algorithm

```
fib( $n$ ) :  
    if  $n \leq 2$  :  $f = 1$   
    else :  $f = \text{fib}(n - 1) + \text{fib}(n - 2)$   
    return  $f$ 
```


- It works. Is it good?



Memorized Recursive Algorithm

```
memo = [ ]  
fib(n) :  
    if  $n$  in memo : return memo[ $n$ ]  
  
    if  $n \leq 2$  :  $f = 1$   
    else :  $f = \text{fib}(n - 1) + \text{fib}(n - 2)$   
  
    memo[ $n$ ] =  $f$   
    return  $f$ 
```

- Benefit?
 - fib(n) only recurses the first time it's called





Dynamic Programming Algorithm

- Memorize (remember) & re-use solutions to subproblems that helps solve the problem
- DP \approx recursion + memorization

Backpropagation

Gradients in ANN

- Learning weights and biases from data using gradient descent
- $\frac{\partial \ell}{\partial \omega}$: too many computations are required for all ω
- Structural constraint of NN:
 - Composition of functions
 - Chain rule
 - Dynamic programming




Training Neural Networks: Backpropagation Learning

- Forward propagation
 - the initial information propagates up to the hidden units at each layer and finally produces output
- Backpropagation
 - allows the information from the cost to flow backwards through the network in order to compute the gradients


Backpropagation

- Chain Rule
 - Computing the derivative of the composition of functions
 - $f(g(x))' = f'(g(x))g'(x)$
 - $\frac{dz}{dx} = \frac{dz}{dy} \cdot \frac{dy}{dx}$
 - $\frac{dz}{dw} = \left(\frac{dz}{dy} \cdot \frac{dy}{dx}\right) \cdot \frac{dx}{dw}$
 - $\frac{dz}{du} = \left(\frac{dz}{dy} \cdot \frac{dy}{dx} \cdot \frac{dx}{dw}\right) \cdot \frac{dw}{du}$
- Backpropagation
 - Update weights recursively




Backpropagation

- Chain Rule
 - Computing the derivative of the composition of functions
 - $f(g(x))' = f'(g(x))g'(x)$
 - $\frac{dz}{dx} = \frac{dz}{dy} \cdot \frac{dy}{dx}$
 - $\frac{dz}{dw} = \left(\frac{dz}{dy} \cdot \frac{dy}{dx}\right) \cdot \frac{dx}{dw}$
 - $\frac{dz}{du} = \left(\frac{dz}{dy} \cdot \frac{dy}{dx} \cdot \frac{dx}{dw}\right) \cdot \frac{dw}{du}$
- Backpropagation
 - Update weights recursively




Backpropagation

- Chain Rule
 - Computing the derivative of the composition of functions
 - $f(g(x))' = f'(g(x))g'(x)$
 - $\frac{dz}{dx} = \frac{dz}{dy} \cdot \frac{dy}{dx}$
 - $\frac{dz}{dw} = \left(\frac{dz}{dy} \cdot \frac{dy}{dx} \right) \cdot \frac{dx}{dw}$
 - $\frac{dz}{du} = \left(\frac{dz}{dy} \cdot \frac{dy}{dx} \cdot \frac{dx}{dw} \right) \cdot \frac{dw}{du}$
- Backpropagation
 - Update weights recursively




Backpropagation

- Chain Rule
 - Computing the derivative of the composition of functions
 - $f(g(x))' = f'(g(x))g'(x)$
 - $\frac{dz}{dx} = \frac{dz}{dy} \cdot \frac{dy}{dx}$
 - $\frac{dz}{dw} = \left(\frac{dz}{dy} \cdot \frac{dy}{dx}\right) \cdot \frac{dx}{dw}$
 - $\frac{dz}{du} = \left(\frac{dz}{dy} \cdot \frac{dy}{dx} \cdot \frac{dx}{dw}\right) \cdot \frac{dw}{du}$
- Backpropagation
 - Update weights recursively with memory



Training Neural Networks with TensorFlow

- Optimization procedure




- It is not easy to numerically compute gradients in network in general.
 - The good news: people have already done all the "hard work" of developing numerical solvers (or libraries)
 - There are a wide range of tools → We will use the TensorFlow

Core Foundation Review


The Perceptron

- Structural building blocks
- Nonlinear activation functions




Neural Networks

- Stacking Perceptrons to form neural networks
- Optimization through backpropagation



Training in Practice

- Adaptive learning
- Batching
- Regularization






(Artificial) Neural Networks with Scikit-learn

**Industrial AI Lab.
Prof. Seungchul Lee**

Our Network Model

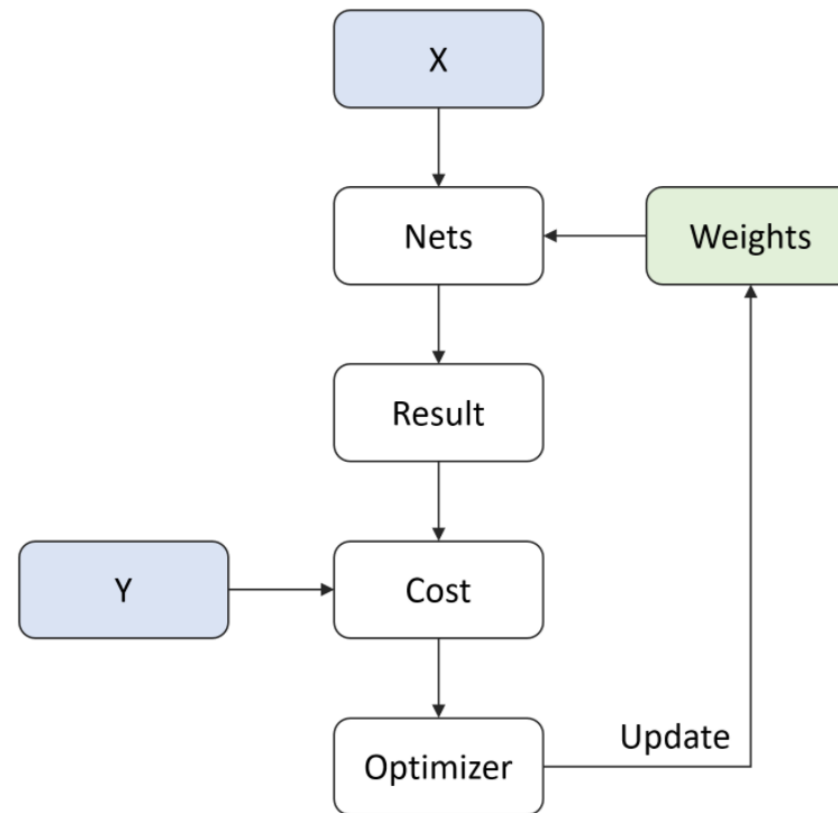


Iterative Optimization

- We will use
 - Mini-batch gradient descent
 - Adam optimizer

$$\begin{array}{ll} \min_{\theta} & f(\theta) \\ \text{subject to} & g_i(\theta) \leq 0 \end{array}$$

$$\theta := \theta - \alpha \nabla_{\theta} \left(h_{\theta} \left(x^{(i)} \right), y^{(i)} \right)$$



ANN with Scikit-learn

- Import Library

```
# Import Library
import numpy as np
import matplotlib.pyplot as plt

from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score
```

- Load MNIST Data

- Download MNIST data


```
train_x = np.load('./data_files/mnist_train_images.npy')
train_y = np.load('./data_files/mnist_train_labels.npy')
test_x = np.load('./data_files/mnist_test_images.npy')
test_y = np.load('./data_files/mnist_test_labels.npy')
```


Training

```
clf = MLPClassifier(hidden_layer_sizes = (100,),  
                    activation = 'logistic',  
                    solver = 'sgd',  
                    learning_rate_init = 0.0001,  
                    batch_size = 50,  
                    max_iter = 100,  
                    verbose = True)
```

```
clf.fit(train_x, train_y)
```

Iteration 1, loss = 3.57824966
Iteration 2, loss = 3.18791200
Iteration 3, loss = 3.13809828
Iteration 4, loss = 3.08519212
Iteration 5, loss = 3.02767935
Iteration 6, loss = 2.96478203
Iteration 7, loss = 2.89614439
Iteration 8, loss = 2.82202289
Iteration 9, loss = 2.74309025
Iteration 10, loss = 2.66058488



Test or Evaluation

```
pred = clf.predict(test_x)
print("Accuracy : {}".format(accuracy_score(test_y, pred)*100))
```

Accuracy : 96.0%

```
logits = clf.predict_proba(test_x[:1])
predict = clf.predict(test_x[:1])

plt.figure(figsize = (6,6))
plt.imshow(test_x[:1].reshape(28,28), 'gray')
plt.xticks([])
plt.yticks([])
plt.show()

print('Prediction : {}'.format(np.argmax(predict)))
np.set_printoptions(precision = 2, suppress = True)
print('Probability : {}'.format(logits.ravel()))
```

Prediction : 7

Probability : [0.02 0. 0.01 0.03 0.01 0.02 0. 0.93 0.01 0.12]

